



Harvesting Tables in the Wild

Jack Henschel, Rohit Raj, Eelis Kostinen


The Knowledge on the Internet

- World Wide Web is vast and has lots of “data”
- Mostly focused on human consumption
- Turning this “data” into knowledge is a semantic challenge!
- *TSLA, 2020-04-06, 420.2*

Our Goal:

**Build a scalable and reliable
pipeline to extract HTML
tables from web pages**

Related Work



Literature Review

- Comprehensive review by S. Zhang and K. Balog^[1] about:
 - a. Table Extraction
 - b. Table Interpretation
 - c. Table Search
 - d. Table Question Answering
 - e. Knowledge Base Augmentation
 - f. Table Augmentation

[1] S. Zhang, K. Balog: “Web Table Extraction, Retrieval, and Augmentation: A Survey” ACM Transactions on Intelligent Systems and Technologies (2020).


Using reference knowledge bases

- Work by Muñoz et. al.^[2] to match tables from Wikipedia
 - Extracts data from wikipedia as RDF
 - Uses the knowledge gained to both augment the knowledge DB
- D. Ritze et. al.^[3] proposed using DBpedia
 - Evaluates performance of HTML to a KB matching system
 - T2K Match: an iterative matching method combining schema and instance matching for matching common HTML tables against cross-domain knowledge bases

[2] E. Muñoz, A. Hogan, A. Mileo: “Using linked data to mine RDF from wikipedia’s tables” WSDM (2014).

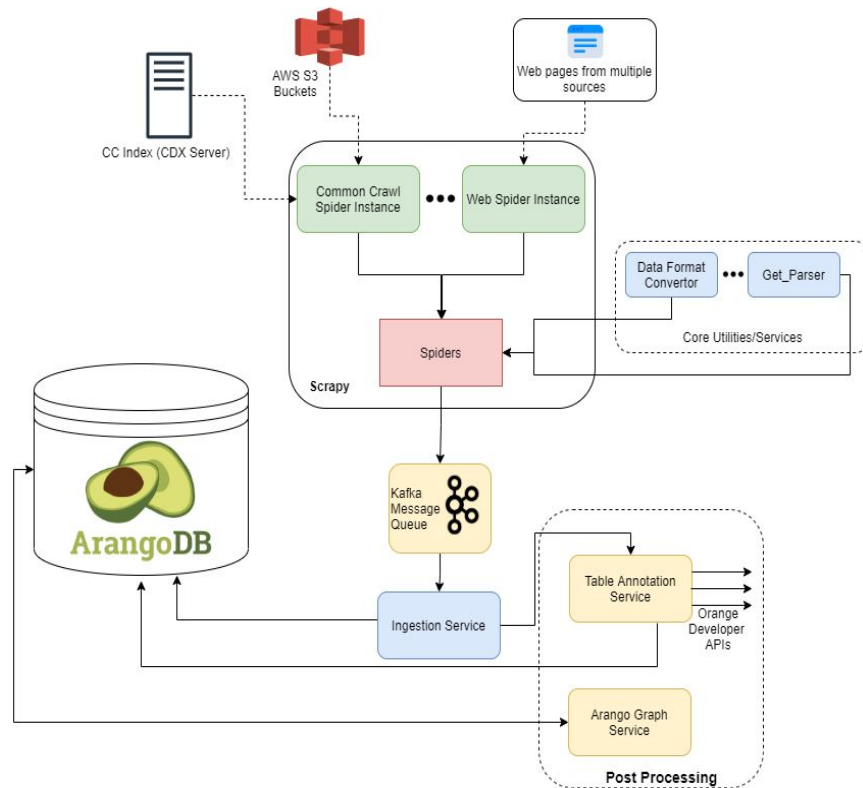
[3] O. Lehmberg, D. Ritze, R. Meusel, C. Bizer: “A Large Public Corpus of Web Tables containing Time and Context Metadata”, WWW 2016.

System requirements & design



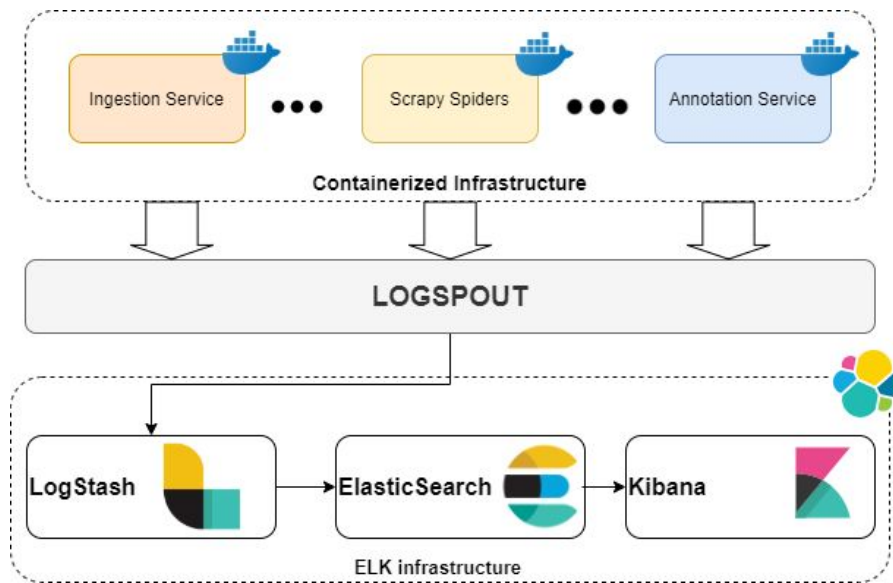
Crawling Pipeline

- High-level requirements
 - Continuous crawling of data from different sources
 - Single-node
- Data sources
 - World Wide Web
 - CommonCrawl (CC) web corpus
- Technologies
 - Crawlers: long-running tasks
 - Storage: represent semantic relationships
 - Ingestion: post-processing scalable to high volumes
 - Monitoring: provide insights to crawling pipeline



Monitoring Infrastructure

- **Problem:** pipeline insights unknown
- **Solution:** continuous monitoring
- Elasticsearch, Logstash, Kibana (ELK stack)
 - Logspout for docker log redirection
 - Explicit redirection



Implementation & Milestones

Getting started

Minimal implementation with Scrapy framework:

1. Download HTML from given URL
2. Extract *<table>* elements
3. Store them in a JSON file

Advanced table parsing

- HTML tables are used for many different purposes
 - need to **select the relevant ones**
- HTML may be **erroneous**
- Some websites have specific patterns for their tables
 - special focus on Wikipedia's ***wikitable***s

Data format

- Retain and store as much information about the web page as possible
- Extensive study about data formats other people have used
- **Final choice:** JSON schema based on DWTC format with additional fields
- Validation through JSON Schema

```
1 {
2   "hasHeader": true,
3   "pageTitle": "COVID-19 situation update for the EU",
4   "url": "https://www.ecdc.europa.eu/en/cases-2019-r",
5   "headerPosition": "FIRST ROW",
6   "tableType": "RELATION",
7   "relation": [
8     [
9       "EU/EEA",
10      "Sum of Cases",
11      "Sum of Deaths",
12      "14-day case notification rate per 100 000 inh",
13      "14-day death notification rate per 1 000 000",
14      "Reporting period YYYY-WW"
15    ],
16    [
17      "France",
18      "3053617",
19      "73049",
20      "403.45",
21      "79.07",
22      "2021-02 and 2021-03"
23    ],
24    [
25      "Spain",
26      "2593382",
27      "56208",
28      "1026.05",
29      "83.79",
30      "2021-02 and 2021-03"
31    ],
32    [
33      "Italy",
34      "2466813",
35      "85461",
36      "315.31"
```

Common Crawl

- Extremely large corpus of web pages (12 years)
- Publicly available index: Data hosted on Amazon S3
- Additional data source for our pipeline
- Data saved in **WARC format**
 - Search the index server
 - Download the WARC file from S3 bucket

Crawling Strategy

- *Which websites should we visit? Where should we start?*
- Crawler follows all links on a web page
 - Two tier regex approach
- **Whitelisted** and **blacklisted domains**
- Custom selection of **Seed URLs**
- Based on Alexa ranking and CC Top 1000

Visualization

Visited web pages (and extracted tables) form a relationship

→ **Graph**

ArangoDB natively supports interacting with graph structures

Vertices: *visited_pages* and *parsed_tables*

Edges: *hyperlink* and *page_contains*

Visualization

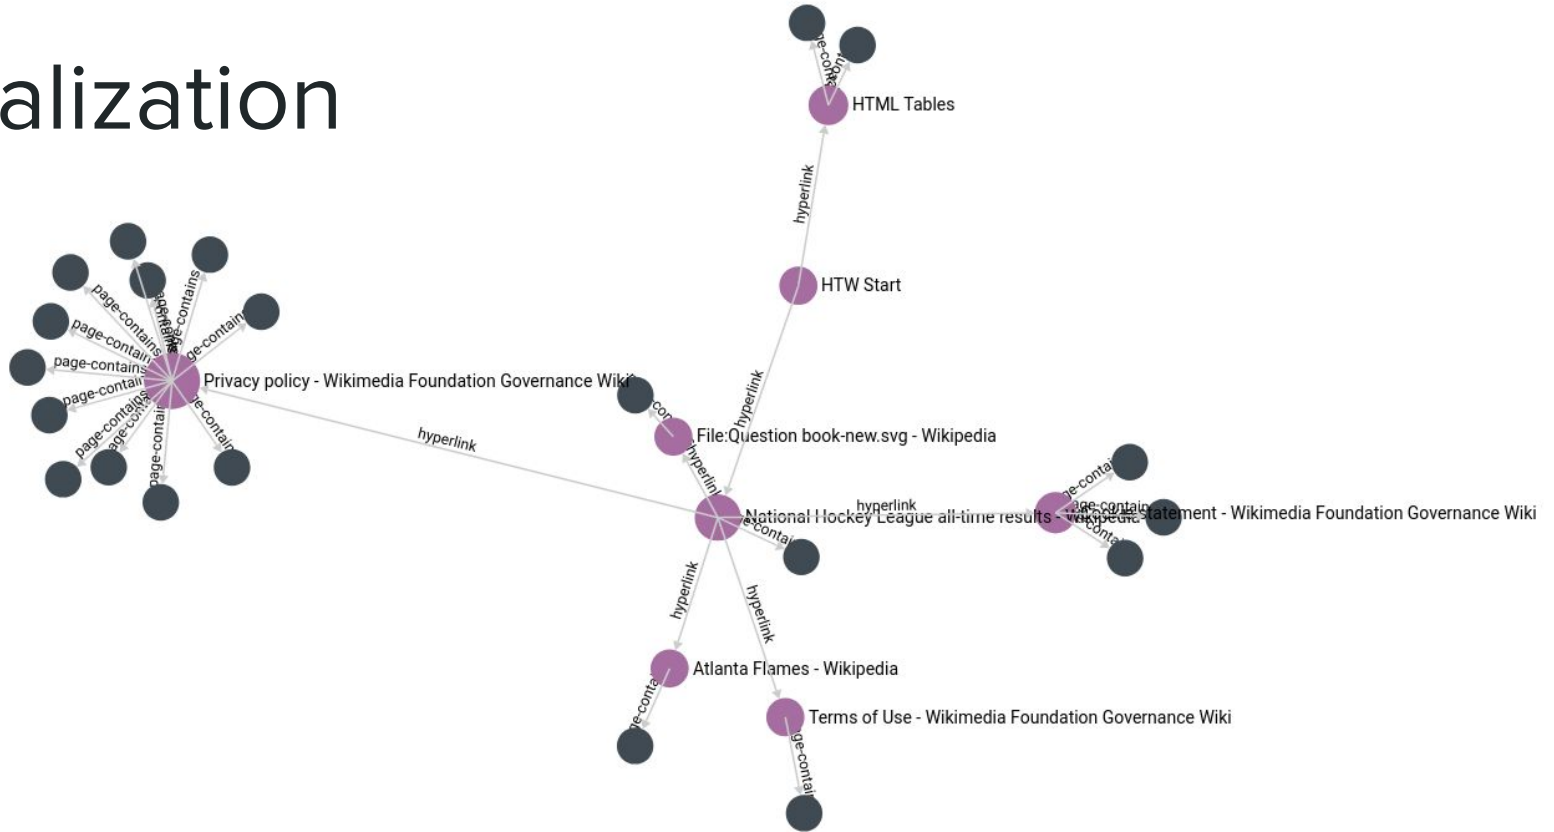
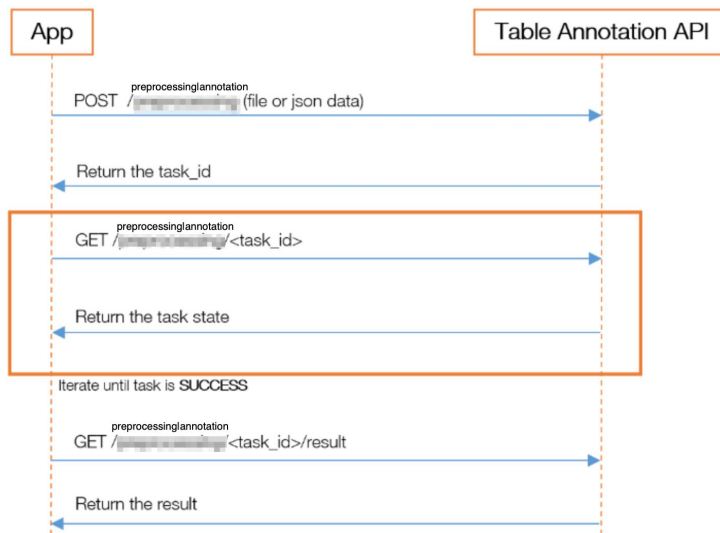


Table Processing

- Goal: extracting meaningful relationships from crawled data
- Solution: integrate pipeline with Orange Table Annotation API ^[4]
- Steps:
 - Preprocessing
 - Annotation



Results



Partial Graph of final collection




Results - Some Stats

- Final run during last week of January
- 3469 tables collected
- 2002 websites visited
- 135 unique domains visited
- Low count due to a constrained whitelist strategy and suboptimal Seed URL selection

Results - Qualitative

- Built a robust crawling pipeline from the ground up
- Used modern technologies
- Based on relevant and recent scientific literature
- Followed software development best practices to keep code extensible and maintainable

Future Work



Pipeline Improvements

- Refresh previously crawled data
- Parse pages with client-side rendering
- Improvements to seed-pages and whitelisted domains
- Optimizations in extraction and pre-processing algorithms

Table Annotation

- Integrating pipeline with annotation service once stable release available
 - Current beta release too buggy for a stable integration
 - `{'code': 1, 'message': 'Internal error', 'description': "apigee - CatchCommonErrors exception: TypeError: Can't use 'in' on a non-object."}`
- Provide more valuable data insights

Thanks for your attention!
Any questions?

Full report available at: <https://u9k.de/htw-report>

